

Cryptographic Hash Functions - Lab Report

CPE-321: Introduction to Computer Security

Team Members:

Abhiram Yakkali

AI Tool Citations:

- Claude (Anthropic) - Used for code generation assistance, debugging, and report writing
- Visual Studio Code - IDE used for development
- Python 3.x with libraries: hashlib, bcrypt, nltk, matplotlib

TASK 1: Exploring Pseudo-Randomness and Collision Resistance

Task 1a: SHA256 Hashing

We built a simple program that takes any input and runs it through SHA256, then displays the resulting hash in hexadecimal format. We used Python's built-in hashlib library for this since it provides a reliable and efficient SHA256 implementation.

Task 1b: Hamming Distance Exploration

To explore how sensitive SHA256 is to input changes, we created pairs of strings that differ by just a single bit and hashed them both. We ran this experiment several times to see how the outputs compared.

Question 1: Observations from Task 1b

What we found: Even when two inputs differ by just a single bit, their SHA256 hashes look completely unrelated. On average, about **128 bits (50%)** of the output changed, and usually **all 32 bytes** ended up being different.

This is called the **avalanche effect**, and it's a key property of good cryptographic hash functions. The idea is that even a tiny change to the input should cause a massive, unpredictable change in the output. This makes it practically impossible to figure out what the original input was just by looking at the hash, or to find patterns that could be exploited.

Task 1c: Finding Collisions

Next, we modified our program to work with truncated hashes (between 8 and 50 bits) so we could actually find collisions in a reasonable amount of time. We used the birthday attack approach, which is much faster than brute force. The trick is to store all the hashes we've computed in a dictionary, so checking for a collision is instant.

Question 2: Collision Analysis

Worst case scenario: For an n -bit hash, you'd need at most $2^n + 1$ attempts to guarantee finding a collision (pigeonhole principle - if there are only 2^n possible outputs, the 2^{n+1} th input must collide with something).

Expected case (Birthday Bound): Thanks to the birthday paradox, we actually expect to find a collision much sooner - around $2^{n/2}$ hashes. This is because collision probability grows quadratically as we add more samples.

What we observed: Our experiments matched the birthday bound predictions really well. The number of hashes needed was consistently close to $2^{n/2}$.

How long would a full 256-bit collision take?

- We'd need around $2^{128} \approx 3.4 \times 10^{38}$ hashes
- At 1 million hashes per second: $\sim 10^{32}$ seconds
- That's roughly **10²⁴ years**
- The universe is only about 1.4×10^{10} years old

So yeah, SHA256 is collision-resistant for any practical purpose.

Question 3: Pre-image Resistance vs Collision Resistance

Can we break the one-way property with an 8-bit digest?

Absolutely. With only 256 possible hash values (2^8), we can easily find an input that produces any given hash - just try random inputs until one works. At most, that's 256 attempts.

Is finding a pre-image easier or harder than finding a collision?

Finding a pre-image is harder. Here's why:

- **Pre-image attack:** You need to find a specific input that produces a specific output. That takes $O(2^n)$ work on average.
- **Collision attack:** You just need any two inputs that hash to the same thing. Thanks to the birthday paradox, that only takes $O(2^{n/2})$ work.

For an 8-bit digest, this means:

- Pre-image: ~ 128 attempts on average
- Collision: ~ 16 attempts (square root of 256)

This is why collision resistance is considered a weaker property than pre-image resistance. Breaking collision resistance is fundamentally easier.

TASK 2: Breaking Real Hashes (Bcrypt)

Implementation Overview

We wrote a custom bcrypt password cracker in Python. The cracker reads the shadow file, pulls out each user's hash and salt, then tries every word from the NLTK dictionary (about 135,000 words between 6-10 characters) until it finds a match.

How it works:

- The shadow file format is: User:\$Algorithm\$Workfactor\$SaltHash
- We extract the salt (first 22 characters after the workfactor) and use it with `bcrypt.checkpw()` to test each guess
- We group users by workfactor so we can process all the fast hashes first
- The program logs progress so we can see how far along it is

Bcrypt Workfactor Analysis

The whole point of bcrypt is to be slow - it's designed to make password cracking painful. The workfactor controls how many iterations it does, and each increment doubles the time:

- Workfactor 8: ~30 ms per hash
- Workfactor 9: ~60 ms per hash
- Workfactor 10: ~110 ms per hash
- Workfactor 11: ~220 ms per hash
- Workfactor 12: ~420 ms per hash
- Workfactor 13: ~840 ms per hash

This exponential scaling is exactly what makes bcrypt effective at slowing down attackers.

Cracking Results

Here's what we found. The time for each password depends on where it appears in the dictionary - words near the beginning get found quickly, while words near the end take much longer.

| User | Workfactor | Password | Time |
|---------|------------|----------|---------------------|
| Bilbo | 8 | welcome | 392.82s (6.5 min) |
| Gandalf | 8 | wizard | 390.51s (6.5 min) |
| Thorin | 8 | diamond | 377.49s (6.3 min) |
| Fili | 9 | desire | 782.39s (13 min) |
| Kili | 9 | ossify | 740.89s (12.3 min) |
| Balin | 10 | hangout | 1557.72s (26 min) |
| Dwalin | 10 | drossy | 1361.81s (22.7 min) |
| Oin | 10 | ispaghul | 1436.17s (24 min) |

| | | | |
|-------|----|------------|---------------------|
| Gloin | 11 | oversave | 2896.63s (48 min) |
| Dori | 11 | indoxylic | 2755.79s (46 min) |
| Nori | 11 | swagsman | 2929.05s (49 min) |
| Ori | 12 | airway | 5450.16s (1.5 hrs) |
| Bifur | 12 | corrosible | 5727.11s (1.6 hrs) |
| Bofur | 12 | libellate | 5837.28s (1.6 hrs) |
| Durin | 13 | purrone | 14509.72s (4.0 hrs) |

Total Cracking Time: 47,145.63 seconds (~13.1 hours)

Method: Parallel dictionary attack using multiprocessing (8 CPU cores)

Dictionary: NLTK word corpus, 6-10 character words (~135,000 words)

Success Rate: 15/15 passwords cracked (100%)

Question 4: Brute Force Time Estimates

Starting point:

- Our dictionary has ~135,000 words
- At workfactor 10, each hash takes about 110 ms
- Worst case for a single word: $135,000 \times 0.11s \approx 4.1$ hours

What about word1:word2 (two dictionary words)?

- Combinations: $135,000^2 = 18.2$ billion possibilities
- Time needed: $18.2 \times 10^9 \times 0.11s = 2.0 \times 10^9$ seconds
- That's about **63 years** of continuous computation

What about word1:word2:word3 (three words)?

- Combinations: $135,000^3 = 2.46 \times 10^{15}$ possibilities
- Time needed: $2.46 \times 10^{15} \times 0.11s = 2.7 \times 10^{14}$ seconds
- That's roughly **8.6 million years**

What about word1:word2:number (with 1-5 digit number)?

- Number options: $10 + 100 + 1000 + 10000 + 100000 = 111,110$
- Total combinations: $135,000^2 \times 111,110 = 2.02 \times 10^{15}$
- Time needed: about **7.0 million years**

Important assumptions:

- Single-threaded, sequential processing
- Worst case (password is the very last one tried)
- Constant hash time (real-world varies a bit)

Even if you threw 1000 CPU cores at this problem, multi-word passwords would still take thousands of years to crack. This really drives home why passphrases (multiple dictionary words strung together) are so much more secure than single-word passwords - each additional word multiplies the search space by 135,000.

CODE APPENDIX

Task 1: SHA256 Implementation (task1_sha256.py)

See the attached file task1_sha256.py for the complete implementation including:

- SHA256 hashing function
- Truncated hash function
- Hamming distance calculation
- Birthday attack collision finder
- Graph generation for collision analysis

Task 2: Bcrypt Cracker (task2_bcrypt_cracker.py)

See the attached file task2_bcrypt_cracker.py for the complete implementation including:

- Shadow file parser
- NLTK word corpus loader
- Password verification using bcrypt.checkpw()
- Workfactor-grouped cracking for efficiency
- Progress reporting and results saving

Interesting Observations

- 1. The avalanche effect is incredibly consistent.** No matter what input you use or which bit you flip, you always end up with roughly 50% of the output bits changing. It's almost eerie how reliable this is.
- 2. The birthday paradox really works.** Our collision experiments matched the theoretical $2^{n/2}$ predictions almost exactly. It's satisfying when the math and the real-world results line up so well.
- 3. Bcrypt's workfactor scaling is predictable.** Every time you bump the workfactor by 1, the time doubles. This makes it easy to plan how long cracking will take.
- 4. Where your password sits in the dictionary matters a lot.** Common words that appear early alphabetically get cracked much faster. This is why "airway" (Ori's password) took less time than expected despite the high workfactor - it's near the start of the dictionary.
- 5. Some users shared the same salt.** We noticed that users with the same workfactor often had identical salts, which means their hashes were probably generated at the same time or with the same parameters. This doesn't make the passwords any easier to crack - it's just an interesting artifact.